

LOW LATENCY C++

C++ ON SEA 2019

DAVID GROSS
WALTHER ZWART

06/02/2019



SYSTEM TUNING

C++ ON SEA 2019



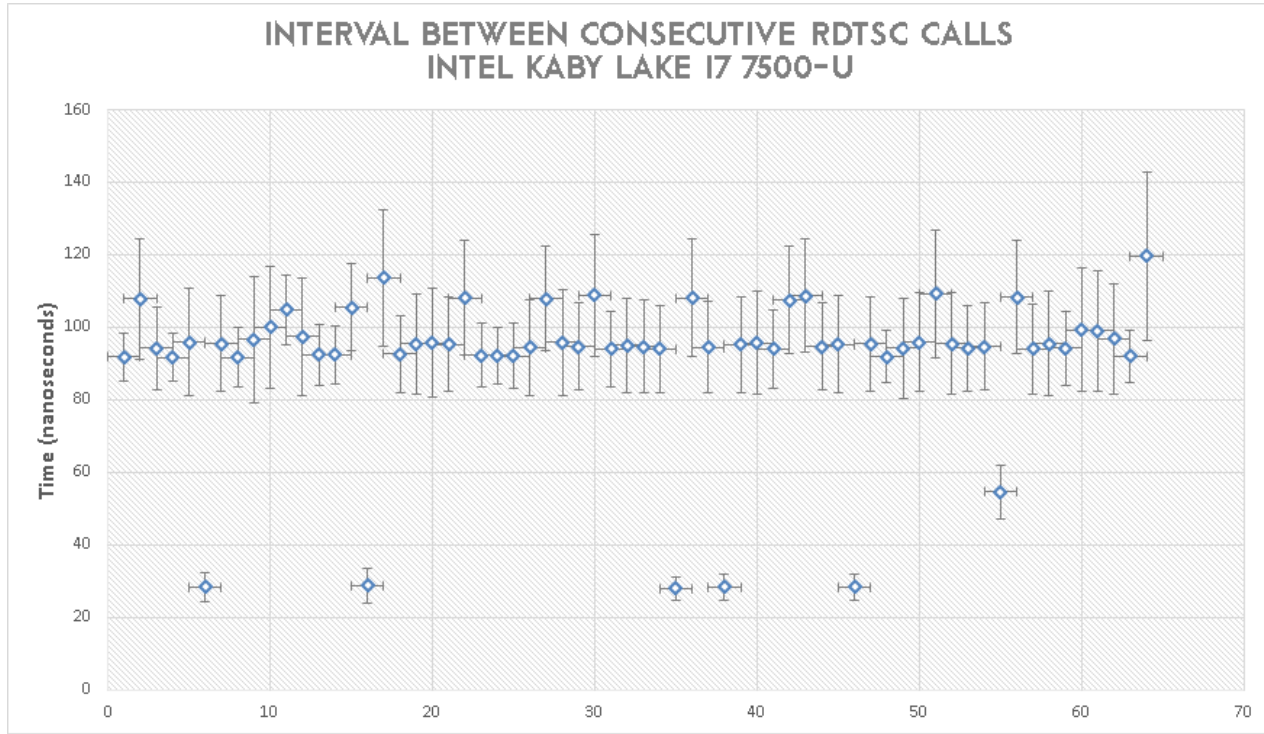
SYSTEM TUNING

- △ Consider the system as a whole
- △ If your workmates write bad code, it will slow down your code as well
- △ Every single part of the system will affect your code performance
- △ It matters especially when you are benchmarking your code!
- △ The best fix is the one that doesn't require writing new code

SYSTEM TUNING

- △ Other applications running on the same server
- △ OS activity
- △ Drivers
- △ Linux kernel configuration
- △ BIOS settings
- △ CPU configuration
- △ Hardware: CPU, RAM, Network card, ...

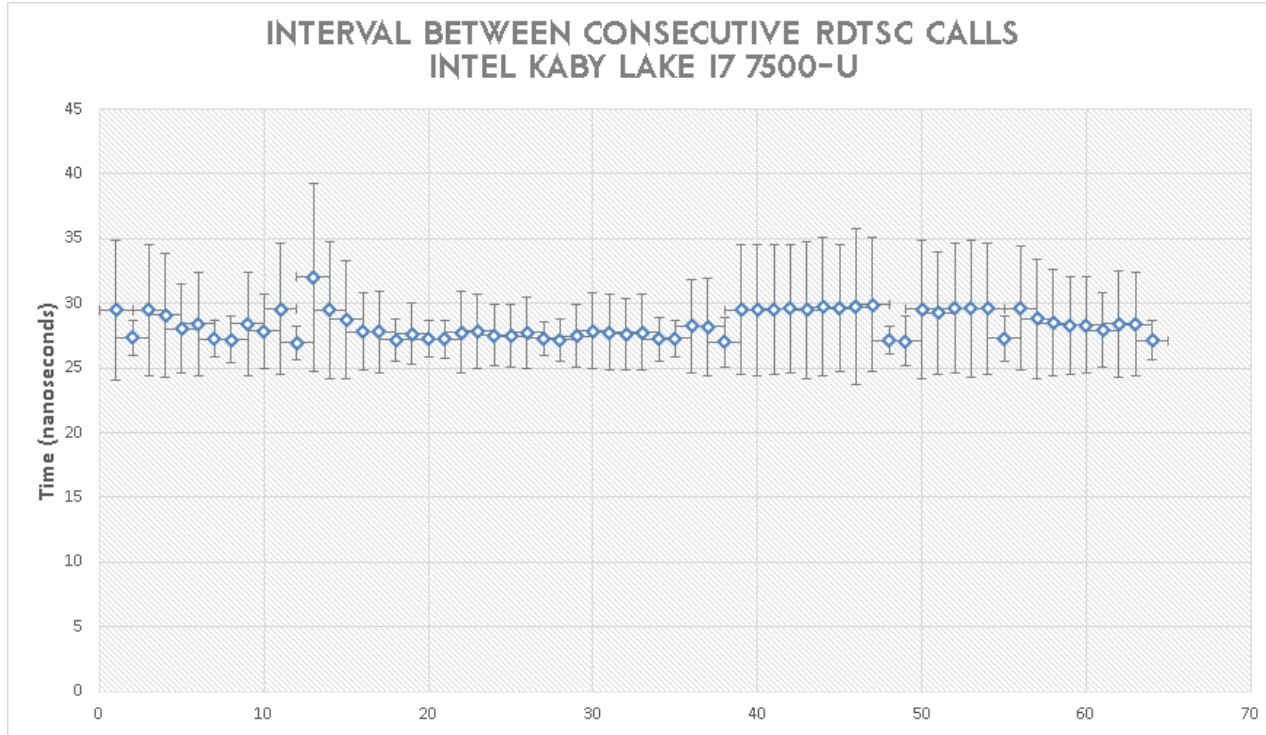
WHEN PREPARING THIS WORKSHOP..



1 sample: 32-256 calls

Sample std dev: 12.23ns

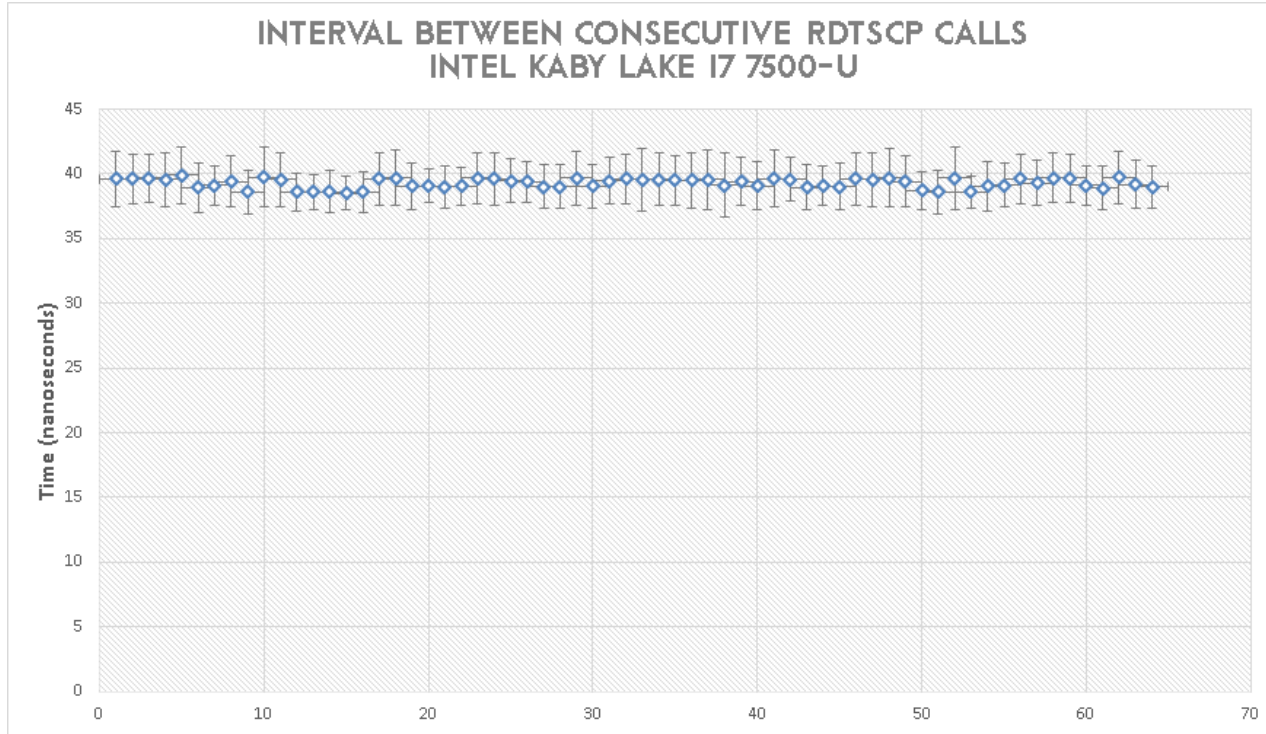
WHEN PREPARING THIS WORKSHOP..



1 sample: 32-256 calls

Sample std dev: 3.49ns

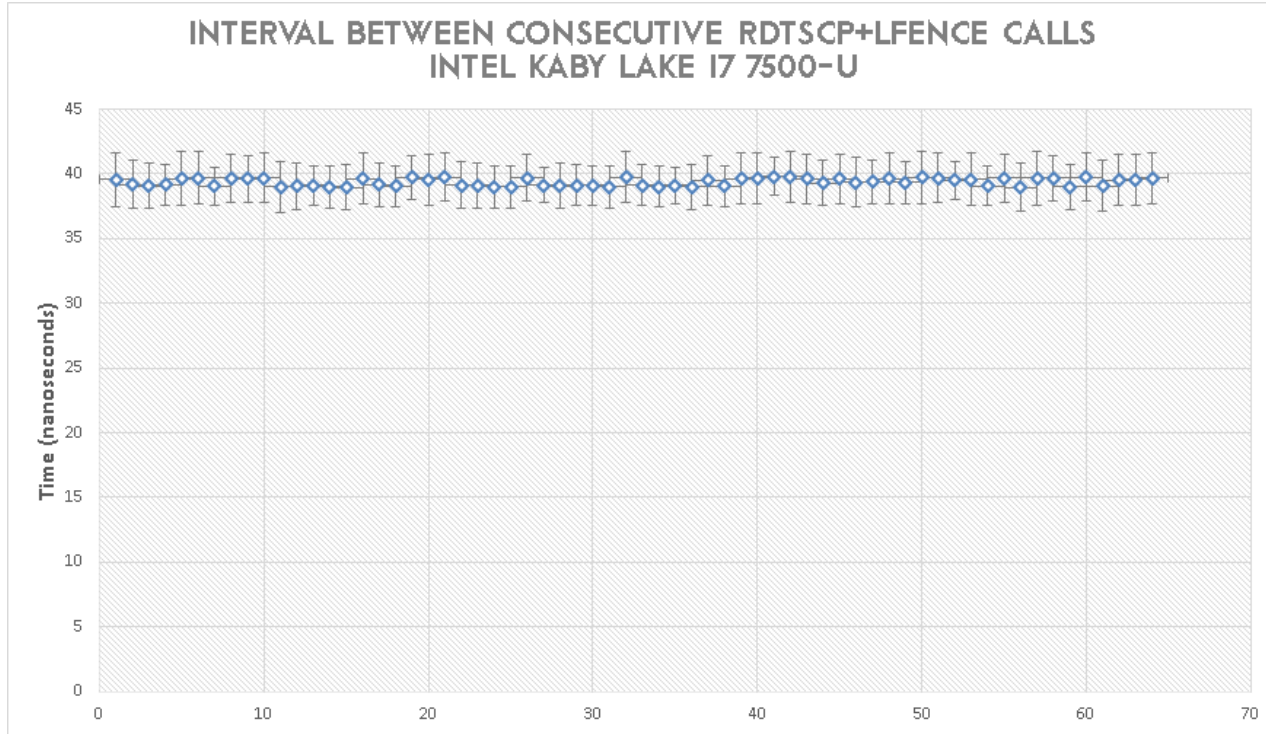
WHEN PREPARING THIS WORKSHOP..



1 sample: 32-256 calls

Sample std dev: 1.85ns

WHEN PREPARING THIS WORKSHOP..



1 sample: 32-256 calls

Sample std dev: 1.72ns

PROCESS AFFINITY

- △ Goal: pin your application to a specific CPU core
- △ Not just your application – all of them
- △ Cache sharing
- △ Interruptions
- △ **Low-latency configuration**
 - △ Run your application via taskset or use sched_setaffinity()

CPU TIME

```
18 [ 0.0%]  
19 [ 0.0%]  
20 [|||||100.0%]  
21 [|||||100.0%]  
22 [ 0.0%]  
23 [ 0.0%]  
24 [|||||100.0%]  
25 [|||||100.0%]  
26 [|||||100.0%]  
Tasks: 66, 50 thr; 14 running  
Load average: 10.54 10.54 10.75  
Uptime: 52 days, 04:43:22
```

P-STATES

- △ Goal: optimize power consumption during code execution
- △ Enhanced Intel Speedstep Technology (EIST) or AMD PowerTune
- △ Each P-State consists of a frequency/voltage pair
- △ Dynamic CPU frequency scaling
- △ **Low-latency configuration**
 - △ Disable EIST and P-States through BIOS
 - △ AND via kernel option `intel_pstate=disable`

C-STATES

△ Goal: optimize power consumption when no code is executed

△ C-states

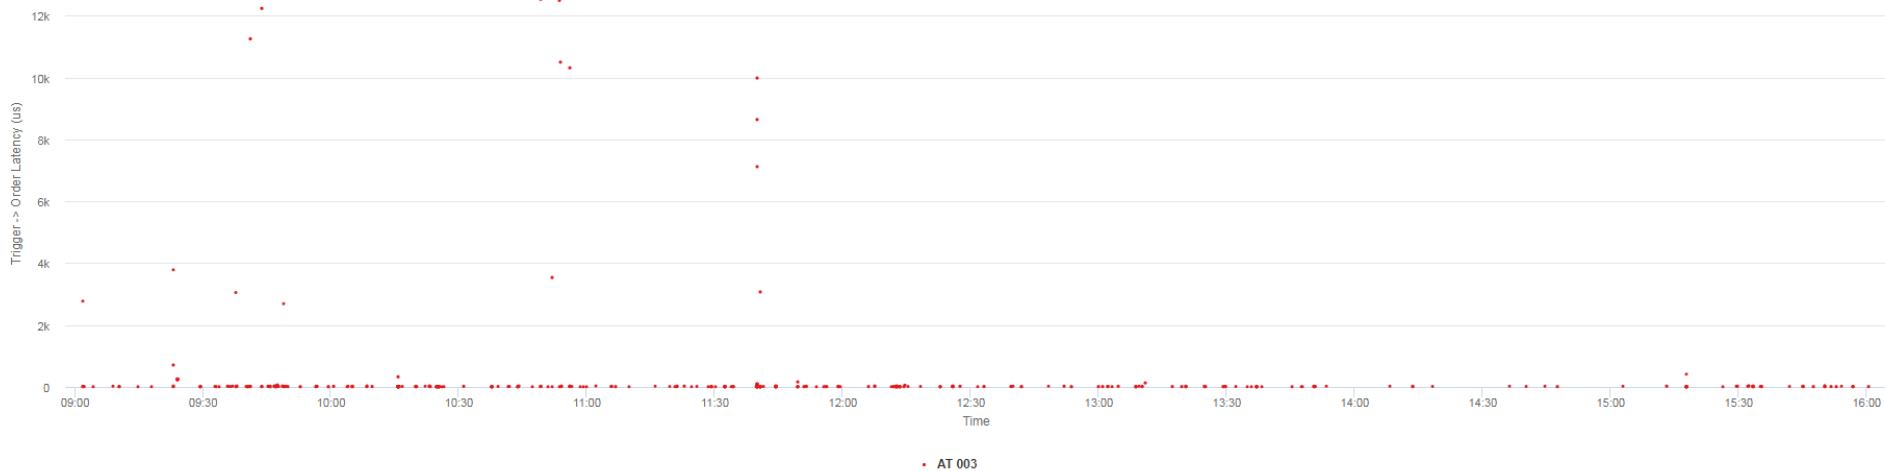
△ C0: Active Mode (code is executed)

△ C3: L1/L2 caches flush

△ C8: LLC flush

C-STATES

Reset Zoom



C-STATES

```
/sys/devices/system/cpu/cpu0/cpuidle$ for i in *; do echo -n "$i: "; cat $i/latency; done  
state1: 2  
state2: 10  
state3: 70  
state4: 85  
state5: 124  
state6: 200  
state7: 480  
state8: 890
```

△ Low-latency configuration

△ Disable C-States through BIOS

△ AND via kernel option `intel_idle.max_cstate=0 idle=poll`

SIMULTANEOUS MULTITHREADING

△ Goal: increase number of instructions executed in parallel

△ Backend stalls

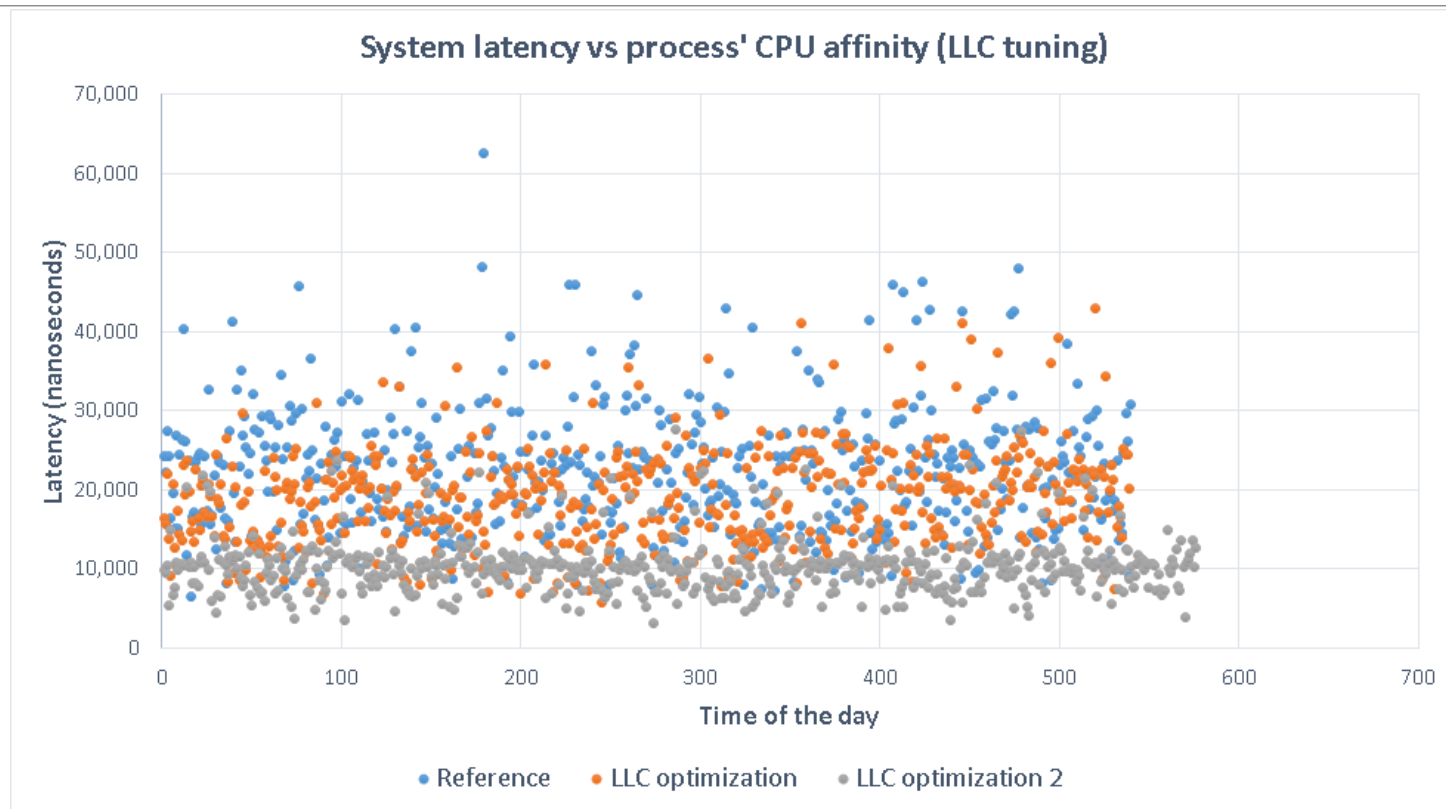
△ Intel Hyper-Threading or AMD SMT

△ Two hardware threads per physical core

△ **Low-latency configuration**

△ Disable SMT through BIOS

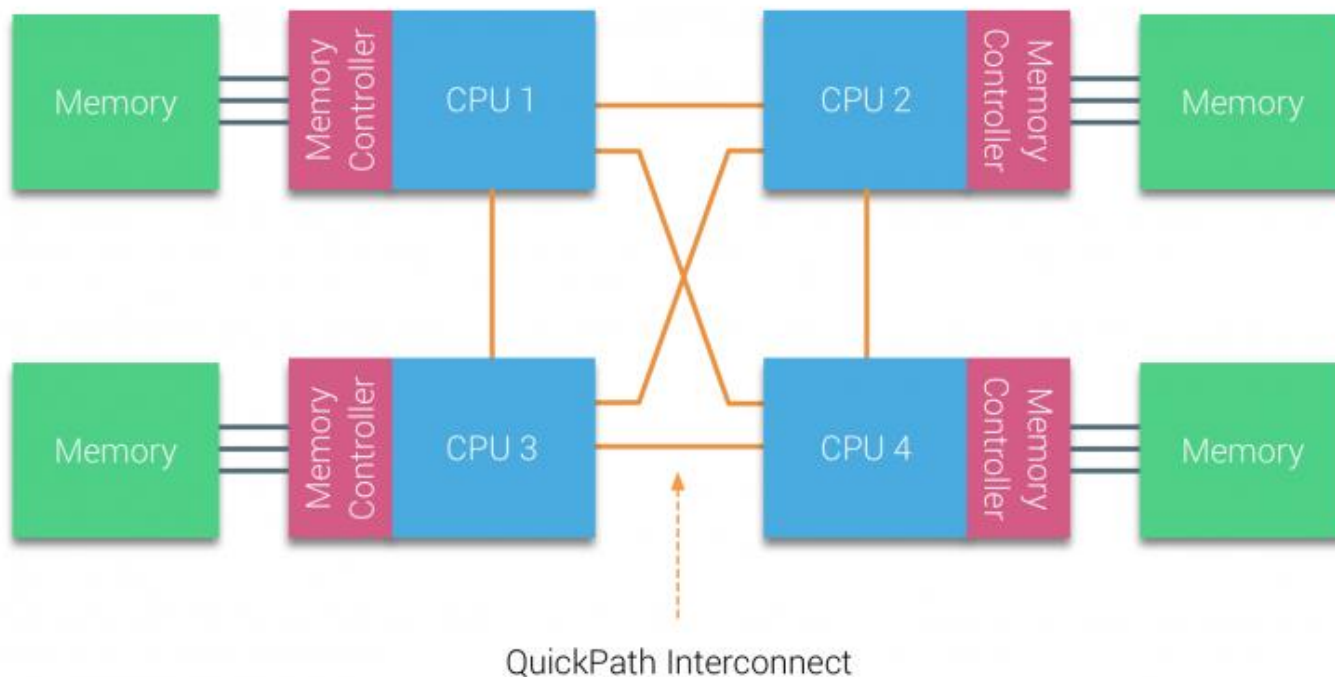
SHARED LLC



SHARED LLC

	Reference	LLC optimization	LLC optimization 2
Min	6,413	5,785	3,092
Median	22,337	19,181	10,181
Max	62,624	42,940	27,654
Std Dev	7,997	6,121	3,493
Mean	22,201	18,997	10,306

NUMA TUNING



Source: frankdenneman.nl

NUMA: ALLOCATION POLICY

- △ Several possible NUMA allocation policies
 - △ Local: allocate memory on the local node
 - △ Preferred: if possible, allocate on X, otherwise on other node
 - △ Membind: forces the allocation on X, if no memory, the process abort
- △ Low-latency configuration
 - △ Run your application via `numactl --membind` or use directly the NUMA API